

API 222 Problem Set 2

Machine Learning and Big Data Analytics: Spring 2024

Oscar Boochever

This problem set is worth 30 points in total. To get full credit, submit your code along with a write-up of your answers. This should either be done in R Markdown or Jupyter Notebook, submitted in one knitted PDF.

Final Project Groups (0 pts)

Please join one of 30 final project groups that have been created on this Canvas page. Details about the final project can be found on this Canvas page. You just need to form your group by March 7. The main project milestones will be after the midterm. We recommend forming groups of 5 students. All students working together should join the same group. PhD students need to work individually (see details on Canvas). If you are a PhD student or are otherwise working alone, please form a group by yourself. Please email Jacob if you have questions about this assignment or the final project.

Conceptual Questions (15 pts)

1. Consider the four main classification methods that have been presented thus far this semester: logistic regression, k-Nearest Neighbors, linear discriminant analysis (LDA), and quadratic discriminant analysis (QDA). Which of these methods may be appropriate if you know the decision boundary between the classes is linear? (3pts)

Logistic regression and LDA.

2. Suppose you had the following data and you are using KNN Regression with Euclidean distance. Consider the prediction problem where you want to predict Y for the data point $X_1 = X_2 = X_3 = 0$.

X1	X2	X3	Y
0	3.5	2	2
1	2.1	3	1
2	4.7	1	3
1	3.9	1	2
0	2.9	2	4
1	1.5	2	1
1	3.5	4	2

- (a) Compute the Euclidean distance between each observation and the test point, $X_1 = X_2 = X_3 = 0$. (1pt)

```

# Make the above data a dataframe
euc_data <- data.frame(
  X1 = c(0, 1, 2, 1, 0, 1, 1),
  X2 = c(3.5, 2.1, 4.7, 3.9, 2.9, 1.5, 3.5),
  X3 = c(2, 3, 1, 1, 2, 2, 4),
  Y = c(2, 1, 3, 2, 4, 1, 2)
)

# Define the x1 = x2 = x3 = 0 as a point
test_point <- c(0, 0, 0)

# Create a function for Euclidean distance
euclidean_distance_knn <- function(x1, x2, x3, test_point) {
  sqrt((x1 - test_point[1])^2 + (x2 - test_point[2])^2 + (x3 - test_point[3])^2)
}

# Apply function to data
euc_data$distance_from_test <- round(euclidean_distance_knn(euc_data$X1, euc_data$X2, euc_data$X3, test_point), 4)

print(euc_data[, c("X1", "X2", "X3", "Y", "distance_from_test")])

```

```

##   X1  X2 X3 Y distance_from_test
## 1  0 3.5 2 2           4.031
## 2  1 2.1 3 1           3.796
## 3  2 4.7 1 3           5.205
## 4  1 3.9 1 2           4.148
## 5  0 2.9 2 4           3.523
## 6  1 1.5 2 1           2.693
## 7  1 3.5 4 2           5.408

```

(b) What is your prediction with $K = 2$? Why? (1pt)

Your prediction is $Y = 2.5$ because the two nearest neighbors (lowest distance values: 3.5 and 2.7 have Y values of 4 and 1 respectively. Thus, the average of these numbers is 2.5.

(c) If the Bayes decision boundary in this problem is highly nonlinear, then would we expect the best value for K to be large or small? Why? (1pt)

We would expect K to be small so the model could be very flexible and sensitive to slight variation and nonlinearity.

- Consider we conduct a research study analyzing the risk factors for developing prostate cancer among men, with variables X_1 = age (years), X_2 = family history of prostate cancer (0 = no, 1 = yes), X_3 = smoking status (0 = non-smoker, 1 = smoker), and Y = probability of developing prostate cancer. A logistic regression analysis is performed, resulting in estimated coefficients $\hat{\beta}_1 = 0.06$, $\hat{\beta}_2 = 1.2$, $\hat{\beta}_3 = 0.8$, and $\hat{\beta}_0 = -3.5$.

(a) Interpret $\hat{\beta}_2$. (1 pt)

$B_2 = 1.2$ means that for men with a family history of prostate cancer, the log-odds of developing prostate cancer are 1.2 units higher relative to those without a family history, on average, holding age and smoking status equal.

- (b) Estimate the probability that a 60-year-old man with a family history of prostate cancer who is a smoker develops prostate cancer. (2 pts)

$$Y = -3.5 + 0.06Age + 1.2FamilyHistory + 0.8Smoker$$

```
age <- 60
family_history <- 1
smoking_status <- 1

# Model
y <- -3.5 + 0.06 * age + 1.2 * family_history + 0.8 * smoking_status

# Exponentiate
probability <- exp(y) / (1 + exp(y))

probability
```

```
## [1] 0.8909032
```

The probability that a 60yr old man with a history of prostate cancer who smokes develops cancer is 89.10%

4. k-fold cross-validation

- (a) Briefly explain how k-fold cross-validation is implemented. (2pts)
- (b) What are the advantages of k-fold cross-validation relative to the validation set approach? (1pt)
5. Suppose you want to minimize the false negative rate in your classification. You run two models: A and B. AUC for Model A is 0.7 and for Model B is 0.8. Can you conclude that you should choose Model B? Why or why not? (3 pts)

Applied Questions (15 pts)

Predicting Hospital Length of Stay

Download data [here](#).

For the next portion of this assignment you will be working with the `LengthOfStay.csv` dataset. This dataset has data points on patients admitted into hospital, indicators of their health condition and how long they were admitted in the hospital.

This is an important problem in healthcare. In order for hospitals to optimize resource allocation, it is important to predict accurately how long a newly admitted patient will stay in the hospital.

1. What are the dimensions of the dataset? (1 pt)
2. Use the `cor()` function to display the correlations of all **continuous** variables in the dataset. Which variables is most highly correlated with `lengthofstay`? (2 pts)

Consider the prediction problem where you want to predict the length of stay for a patient (`lengthofstay`) against all other variables available in the data set.

3. Run ridge regression with cross-validation and standardized features using the canned function `cv.glmnet` from the package `glmnet`. You can use the λ sequence generated by `cv.glmnet` (you do not need to provide your own λ sequence). In order to receive credit for this question, make the line immediately preceding this command say `set.seed(222)` and run the two lines together. Please report all numbers by rounding to three decimal places. (2 pts)
- (a) Which λ had the lowest mean cross-validation error? (1 pt)
 - (b) What was the cross-validation error? (1 pt)
 - (c) What was the standard error of the mean cross-validation error for this value of λ ? (1 pt)
 - (d) What was the largest value of λ whose mean cross validation error was within one standard deviation of the lowest cross-validation error? (1 pt)
5. Now consider the same prediction problem. Implement your own 5-fold cross-validation routine for KNN for $K = 1, \dots, 50$ (write the cross-validation routine yourself rather than using a canned package). Include the snippet of code you wrote here. It should not exceed 20 lines. (6pts)
- (a) Plot of mean cross-validation MSE as a function of k .
 - (b) The best k according to CV is
 - (c) The cross-validation error for the best k is